

# llamafactory多机多卡

## 1:gpu

gpu使用情况

许伯阳: 5, 6, 7

康总: 3, 4

叶哥: 0

多机多卡验证

master:1,

node1: 2,

## 2:构建镜像,配置

```
docker build -f ./docker/docker-cuda/Dockerfile \
  --build-arg INSTALL_BNB=true \
  --build-arg INSTALL_VLLM=true \
  --build-arg INSTALL_DEEPSPEED=true \
  --build-arg INSTALL_FLASHATTN=true \
  --build-arg PIP_INDEX=https://pypi.org/simple \
  -t llamafactory_mgpus:latest .

docker run -dit --gpus=all \
  -v ./hf_cache:/root/.cache/huggingface \
  -v ./ms_cache:/root/.cache/modelscope \
  -v ./data:/app/data \
  -v ./output:/app/output \
  -p 7860:7860 \
  -p 8000:8000 \
  --shm-size 16G \
  --name llamamaster \
  llamafactory_mgpus
```

配置master和node1

```
docker run -dit --gpus '"device=1"' \
  -v /mnt/llamafactory/mgpus/master/huggingface:/root/.cache/huggingface \
  -v /mnt/llamafactory/mgpus/master/modelscope:/root/.cache/modelscope \
  -v /mnt/llamafactory/mgpus/master/data:/app/data \
  -v /mnt/llamafactory/mgpus/master/output:/app/output \
  -p 4007:7860 \
  -p 4008:8000 \
  -p 4009:9000 \
  --shm-size 32G \
  --name llamamaster \
  llamafactory_mgpus

docker exec -it llamamaster bash
# master
172.17.0.12      de237c11dd3f
# node1
172.17.0.13      5fb58c8c9bf0
```

```
ssh-rsa
AAAAB3NzaC1yc2EAAAQABAAQgQDnB4IZvbwz7HczEDP7kwUeYoIRUxbTTenSzex+G1c7VKZ7tuHpG
TSccYwEyy9UuZlmD68MM5zJ4jdUDxfztwQbzVs2w4F4Bxbj125F2pgor4XfSXmQjp56QrKxPHxoLPYI/
2DbB2stA27CdR70zKix8BtBhyQ5WAr9rMhiofU2xRnmZXP7GdJat0/WGxSgTfFaDltVQcMSPy20hsgbtF
zqpXwIIK4ZQdD4hPNooLW+qr2smWDjosCa6NuAqzaTePpbYwIWKfjRrw941j0M3U1Myx9bBDIUrH4vowS
gSsLB0q2AkKc5CcEJdIWNEi17IKIoBZJdQqiuv3nwkhj60n6Bvg8T97n5T689K5B5yjBlhJIGbhTEK1B
GQKjVh+OG1o2bd416MR/viEaiul4ezIbB/6NBqLeGdVJmwu2kAePb7qRgt5TdgoDv2DSJcKa+WUDxmwfZ
UPdt8oaVQ3h0y3bn95fFYAKEIS6GhDkSvZANeoV0Tfe372QZT8dhF0Ksc= root@de237c11dd3f
```

```
docker run -dit --gpus '"device=2"' \
-v /mnt/llamafactory/mgpus/node1/huggingface:/root/.cache/huggingface \
-v /mnt/llamafactory/mgpus/node1/modelscope:/root/.cache/modelscope \
-v /mnt/llamafactory/mgpus/node1/data:/app/data \
-v /mnt/llamafactory/mgpus/node1/output:/app/output \
-p 4010:7860 \
-p 4011:8000 \
-p 4012:9000 \
--shm-size 32G \
--name llamanode1 \
llamafactory_mgpus
docker exec -it llamanode1 bash
```

```
# etc/hosts
# master
172.17.0.12      de237c11dd3f
# node1
172.17.0.13      5fb58c8c9bf0
```

```
ssh-rsa
AAAAB3NzaC1yc2EAAAQABAAQgQDYsdT8FERQbLnbczeASvkzuLk28xtYCgwJ6iHrbHw0xu1G60ydx
vcfERs l+sjKCfan44vWRw8RES00V/Kdfu/+MUV1enmYw4F26V06Rz0o/1bPGJ7PnxaiI+3UNT3X0693ij
66XNQozZ3ZVMVGwuPMeil0lTnKZjfMbyDWQSbmNfMmxmY9KsaogBN37AQE8jX0YxDy0H5aVsYxnByL/xZ
7yVMHvbNHiNctkgGENC2H80R3d57F/Y5d0t8AcM9N1iE/85paX6z/06LIUVzPSLE98f3LdljWL3+5YndA
ixMwXw06S5Rhmsb0MWj4wKmps1QqWNsxLwTPeKcf1prqtPtcrVNjC0Akm87qTrboLsDxl97U0uba0Ts0s
6ad5dZZs/c2PeM69iAZbpvmDeHb4ZFwoTIq5hmMx5CyHbfmM9PcZULyHvfsfo6ccYN4PZ3THARedKh4xP
MwlFF/PSV1YFvFCCS24mzwWR09TwmN2xcUKFCg2snuEwVDg9aBbmaJe1k= root@5fb58c8c9bf0
```

```
# 容器内
chmod 777 /tmp
apt update
apt install -y iputils-ping openssh-server net-tools pdsh
/etc/init.d/ssh start

# hostfile, slots代表使用1个GPU
de237c11dd3f slots=1
5fb58c8c9bf0 slots=1
```

```
~/.ssh/authorized_keys
```

```
ssh-rsa
AAAAB3NzaC1yc2EAAAQABAAQgQDnB4IZvbwz7HczEDP7kwUeYoIRUxbTTenSzex+G1c7VKZ7tuHpG
TSccYwEyy9UuZlmD68MM5zJ4jdUDxfztwdQbzVs2w4F4Bxbj125F2pgor4XfSXmQJp56QrKxPHxoLPYI/
2DbB2stA27CdR70zKix8BtBhyQ5WAr9rMhiofU2xRnmZXP7GdJat0/WGxSgTffFaDltVQcMSPy20hsgbtF
zqpXwIIK4ZQdD4hPNooLW+qr2smWDjosCa6NuAqzaTePpbYwIWkfjRrw941j0M3U1Myx9bBDIUrH4vowS
gSsLB0q2AkKc5CcEJdIWNFEil7IKIoBZJdQqiuV3nwkhj60n6Bvg8T97n5T689K5B5yjBlhJIGbhTEK1B
GQKjVh+OG1o2bd416MR/viEaiul4ezIB/6NBqLeGdVJmwu2kAePb7qRgt5TdgoDv2DSJcKa+WUDxmWFZ
UPdt8oaVQ3h0y3bn95fFYAKEIS6GhDkSvZANeoV0Tfe372QZT8dhF0Ksc= root@de237c11dd3f
```

```
ssh-rsa
AAAAB3NzaC1yc2EAAAQABAAQgQDYsdT8FERQbLnbczeASvkzuLk28xtYCgwJ6iHrbHw0xu1G60ydX
vcfERsl+sjKCfan44vWRw8RES00V/Kdfu/+MUV1enmYW4F26V06RzOo/1bPGJ7PnxaiI+3UNT3X0693ij
66XNQozZ3ZVMVGwuPMef0lTnKZjfMbyDWQSbmNfMmxmY9KsaogBN37AQE8jX0YxDy0H5aVsYxnByL/xZ
7yVMHvbNHiNCtkgGENC2H80R3d57F/Y5d0t8AcM9N1iE/85paX6z/06LIUVzPSLE98f3LdljWL3+5YndA
ixMwXw06S5RhmSb0MWj4wKmps1QqWNsxLwTPeKcf1prqtPtcrVNjC0Akm87qTrboLsDxl97U0uba0Ts0s
6ad5dZZs/c2PeM69iAZbpvmDeHb4ZFwoTIq5hmMx5CyHbfmN9PcZULyHvfsfo6ccYN4PZ3THARedKh4xP
MwlFF/PSV1YFvFCCS24mzWWR09TwmN2xcUKFCg2snuEwVDg9aBbmaJe1k= root@5fb58c8c9bf0
```

ds\_config.json

```
{
  "train_batch_size": "auto",
  "train_micro_batch_size_per_gpu": "auto",
  "gradient_accumulation_steps": "auto",
  "gradient_clipping": "auto",
  "zero_allow_untested_optimizer": true,
  "fp16": {
    "enabled": "auto",
    "loss_scale": 0,
    "initial_scale_power": 16,
    "loss_scale_window": 1000,
    "hysteresis": 2,
    "min_loss_scale": 1
  },
  "zero_optimization": {
    "stage": 2,
    "allgather_partitions": true,
    "allgather_bucket_size": 5e8,
    "reduce_scatter": true,
    "reduce_bucket_size": 5e8,
    "overlap_comm": false,
    "contiguous_gradients": true
  }
}
```

多机sft

```
deepspeed --hostfile /app/config_file/hostfile \
/app/src/train.py --deepspeed /app/config_file/ds_config.json \
--stage sft \
--model_name_or_path /app/output/models/qwen1.5_0.5b_chat \
--do_train True \
```

```
--dataset_dir /app/data/ \
--dataset xhs \
--template qwen \
--finetuning_type lora \
--output_dir /app/output/lora/2 \
--overwrite_cache \
--overwrite_output_dir \
--per_device_train_batch_size 4 \
--gradient_accumulation_steps 4 \
--lr_scheduler_type cosine \
--logging_steps 10 \
--save_steps 10000 \
--learning_rate 1e-5 \
--num_train_epochs 20 \
--plot_loss \
--fp16
```

正常加载:

```
de237c11dd3f: [INFO|trainer.py:2313] 2024-11-20 12:18:50,086 >> ***** Running training *****
de237c11dd3f: [INFO|trainer.py:2314] 2024-11-20 12:18:50,086 >> Num examples = 1,508
de237c11dd3f: [INFO|trainer.py:2315] 2024-11-20 12:18:50,086 >> Num Epochs = 20
de237c11dd3f: [INFO|trainer.py:2316] 2024-11-20 12:18:50,086 >> Instantaneous batch size per device = 4
de237c11dd3f: [INFO|trainer.py:2319] 2024-11-20 12:18:50,087 >> Total train batch size (w. parallel, distributed & accumulation) = 32
de237c11dd3f: [INFO|trainer.py:2320] 2024-11-20 12:18:50,087 >> Gradient Accumulation steps = 4
de237c11dd3f: [INFO|trainer.py:2321] 2024-11-20 12:18:50,087 >> Total optimization steps = 940
de237c11dd3f: [INFO|trainer.py:2322] 2024-11-20 12:18:50,091 >> Number of trainable parameters = 3,784,704
  0% | 0/940 [00:00:<?, ?t/s]/usr/local/lib/python3.10/dist-packages/torch/utils/checkpoint.py:295: FutureWarning: `torch.cpu.amp.autocast(args` is deprecated. Please use `torch.amp.autocast('cpu', args...)` instead.
de237c11dd3f: with torch.enable_grad(), device_autocast(ctx, torch.cpu.amp.autocast(**ctx.cpu_autocast_kwargs)): # type: ignore[attr-defined]
de237c11dd3f: {'loss': 3.9087, 'grad_norm': 0.7430564761161804, 'learning_rate': 9.997207818651273e-06, 'epoch': 0.21}
de237c11dd3f: {'loss': 3.8446, 'grad_norm': 0.6829023957252502, 'learning_rate': 9.988834393115768e-06, 'epoch': 0.42}
de237c11dd3f: {'loss': 3.7103, 'grad_norm': 0.6964868307113647, 'learning_rate': 9.97488907544252e-06, 'epoch': 0.63}
de237c11dd3f: {'loss': 3.7498, 'grad_norm': 0.4726545214653015, 'learning_rate': 9.955387440773902e-06, 'epoch': 0.85}
de237c11dd3f: {'loss': 4.0669, 'grad_norm': 0.48675575852394104, 'learning_rate': 9.930351269950144e-06, 'epoch': 1.06}
de237c11dd3f: {'loss': 3.7059, 'grad_norm': 0.5517176985740662, 'learning_rate': 9.899808525182935e-06, 'epoch': 1.27}
de237c11dd3f: {'loss': 3.6769, 'grad_norm': 0.38814425468444824, 'learning_rate': 9.863793318825186e-06, 'epoch': 1.48}
de237c11dd3f: {'loss': 3.6424, 'grad_norm': 0.3534882664680481, 'learning_rate': 9.822345875271884e-06, 'epoch': 1.69}
de237c11dd3f: {'loss': 3.7453, 'grad_norm': 0.3888565003871918, 'learning_rate': 9.775512486034564e-06, 'epoch': 1.9}
de237c11dd3f: {'loss': 3.9357, 'grad_norm': 0.4186440408229828, 'learning_rate': 9.723345458039595e-06, 'epoch': 2.12}
de237c11dd3f: {'loss': 3.6525, 'grad_norm': 0.3147875964641571, 'learning_rate': 9.665903055208013e-06, 'epoch': 2.33}
de237c11dd3f: {'loss': 3.7402, 'grad_norm': 0.3422832190998448, 'learning_rate': 9.603249433382145e-06, 'epoch': 2.54}
de237c11dd3f: {'loss': 3.5721, 'grad_norm': 0.33048951625823975, 'learning_rate': 9.535454568671705e-06, 'epoch': 2.75}
```

两个GPU占用

master: 1,  
node1: 2,

localhost.localdomain	Wed Nov 20 20:21:50 2024	550.127.05
[0] NVIDIA GeForce RTX 4090	29°C, 0 %   496 / 24564 MB	root(490M)
[1] NVIDIA GeForce RTX 4090	56°C, 52 %   19568 / 24564 MB	root(19562M)
[2] NVIDIA GeForce RTX 4090	59°C, 79 %   20842 / 24564 MB	root(20836M)
[3] NVIDIA GeForce RTX 4090	29°C, 0 %   3 / 24564 MB	
[4] NVIDIA GeForce RTX 4090	28°C, 0 %   3 / 24564 MB	
[5] NVIDIA GeForce RTX 4090	31°C, 0 %   3 / 24564 MB	
[6] NVIDIA GeForce RTX 4090	29°C, 0 %   3 / 24564 MB	
[7] NVIDIA GeForce RTX 4090	31°C, 0 %   3 / 24564 MB	

结束, master 节点文件如下

```

de237c11dd3f: ***** train metrics *****
de237c11dd3f: epoch = 19.8942
de237c11dd3f: total_flos = 33548999GF
de237c11dd3f: train_loss = 3.5809
de237c11dd3f: train_runtime = 0:14:49.74
de237c11dd3f: train_samples_per_second = 33.897
de237c11dd3f: train_steps_per_second = 1.056
de237c11dd3f: Figure saved at: /app/output/lora/2/training_loss.png
de237c11dd3f: [WARNING|2024-11-20 12:33:40] llamafactory.extras.plotting:162 >> No metric eval_loss to plot.
de237c11dd3f: [WARNING|2024-11-20 12:33:40] llamafactory.extras.plotting:162 >> No metric eval_accuracy to plot.
de237c11dd3f: [INFO|modelcard.py:449] 2024-11-20 12:33:40.593 >> Dropping the following result as it does not have all the necessary fields:
de237c11dd3f: {task: {'name': 'Causal Language Modeling', 'type': 'text-generation'}}
5fb58c8c9bf0: [rank1]:[W1120 12:33:40.568127662 ProcessGroupNCCL.cpp:1168] Warning: WARNING: process group has NOT been destroyed before we destruct ProcessGroupNCCL. On normal program exit, the application should call destroy_process_group to ensure that any pending NCCL operations have finished in this process. In rare cases this process can exit before this point and block the progress of another member of the process group. This constraint has always been present, but this warning has only been added since PyTorch 2.4 (function operator())
de237c11dd3f: [rank0]:[W1120 12:33:41.802962566 ProcessGroupNCCL.cpp:1168] Warning: WARNING: process group has NOT been destroyed before we destruct ProcessGroupNCCL. On normal program exit, the application should call destroy_process_group to ensure that any pending NCCL operations have finished in this process. In rare cases this process can exit before this point and block the progress of another member of the process group. This constraint has always been present, but this warning has only been added since PyTorch 2.4 (function operator())
de237c11dd3f: [2024-11-20 12:33:42,465] [INFO] [launch.py:351:main] Process 14741 exits successfully.
5fb58c8c9bf0: [2024-11-20 12:33:42,679] [INFO] [launch.py:351:main] Process 9844 exits successfully.
root@de237c11dd3f:/app# ls /app/output/lora/2/
README.md      added_tokens.json    merges.txt          tokenizer.json   trainer_log.jsonl  training_loss.png
adapter_config.json  all_results.json  runs             tokenizer_config.json  trainer_state.json  vocab.json
adapter_model.safetensors checkpoint-940  special_tokens_map.json  train_results.json  training_args.bin

```

node1的输出如下

```

root@5fb58c8c9bf0:/app/output/lora/2# tree
.
└── checkpoint-940
    ├── global_step940
    │   └── zero_pp_rank_1_mp_rank_0_optim_states.pt
    └── rng_state_1.pth

2 directories, 2 files

```

(master 中执行)合并结果

```

llamafactory-cli export \
--model_name_or_path /app/output/models/qwen1.5_0.5b_chat \
--adapter_name_or_path /app/output/lora/2 \
--template qwen \
--finetuning_type lora \
--export_dir /app/output/exportout/mgp_qwen1.5_chat0.5b \
--export_size 2 \
--export_legacy_format False

```

```

root@de237c11dd3f:/app# ls /app/output/exportout/mgp_qwen1.5_chat0.5b/
added_tokens.json generation_config.json model.safetensors tokenizer.json vocab.json
config.json merges.txt special_tokens_map.json tokenizer_config.json
root@de237c11dd3f:/app#

```

(master 中执行)推理

```

llamafactory-cli chat \
--model_name_or_path /app/output/exportout/mgp_qwen1.5_chat0.5b \
--template qwen

```

```

[INFO|2024-11-20 12:47:39] llamafactory.model.model_utils.attention:157 >> Using torch SDPA for faster training and inference.
[INFO|2024-11-20 12:47:39] llamafactory.model.loader:157 >> all params: 463,987,712
Welcome to the CLI application, use `clear` to remove the history, use `exit` to exit the application.

User: 你好？你能写一个关于提升英语的小红书文案吗？
Assistant: ✨【新学的英语新技能】
★【提升英语能力的12个秘诀】
每天练习单词！
多听多说！
多读多写！
多说多练！
多听多记！

```